A Bottom-Up Approach to Job Recommendation System

(RecSys Competition 2016)

Stoyan Eftimov
Hrizantema Stancheva

Problem

Recommender Systems held during 15-19th September 2016, Boston @ MIT & IBM

RecSys Challenge:

- Build a job recommendation system for XING
- Given a XING user, the goal is to predict those job postings that a user will positively interact with (e.g. click, bookmark)

XING

Find jobs, contacts, events...

Advanced Search

XING

Fabian Abel

PREMIUM

- My start page
- My contacts 25
- My messages 51
- My Premium
  NEW: XING Arbeitsrechtsschutz
- Jobs
- Events

Network news | Comments and likes

Jobs we think you'll like

DevOps Engineer (m/f) for Data...
XING AG

Software Architekt (m/w) mit d...
adesso AG

Projektleiter (m/w) im Bereich...
adesso AG

(Senior) Consultant Data Wareh...
empiricus GmbH - Agentur für I...

> 16 more job recommendations

Share something with your contacts

What's new?

Datasets

Users	Items	Interactions
Job roles	Title	User_ID
Career level	Discipline	Item_ID
Discipline	Industry	Time
Industry	Country	Interaction_type
Country	Region	
Region	Type of employment	
Work experience	Tags	
Education	Creation time	

Interactions Data

- Purely focus on interactions

1 - user clicked on the item

2 - user bookmarked the item on XING

3 - user clicked the reply button or application form button

4 - user deleted the recommendation

- Issues:

Data is sparse and doesn’t include all users

Not every user has at least 30 positive interactions

Collaborative Filtering

- Leverages the notion

User-user similarity

Item-item similarity

- Challenges:

Sparsity

Definition of Similarity

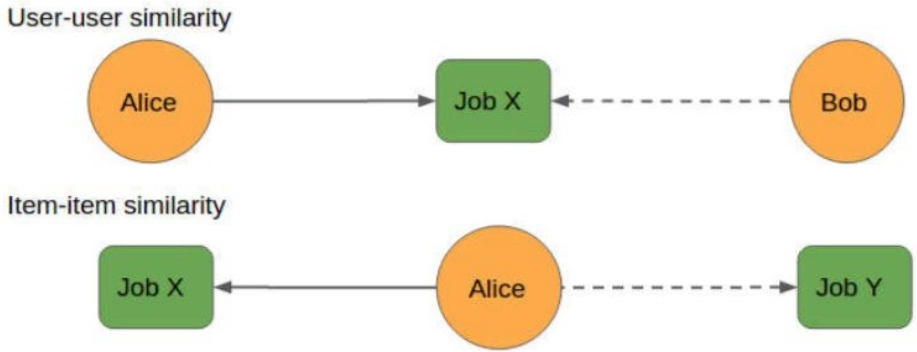

Collaborative Filtering

- For similarity, we use K-Means clustering with cosine similarity
- Impute the interaction of user u with item i as average of interactions of other users v, in the same cluster, weighted by cosine similarity

Collaborative Filtering

How to we know if two items or users are similar?

K-means Clustering

- Number of clusters: 100, 1K, 5K
- Distance measure : Euclidean distance
- Library: SciPy.kmeans2

Collaborative Filtering

Cosine Similarity

- Similarity (U i , U j) = U i .U j / |U i | |U j |
- Similarity (I i , I j) = I i .I j / |I i | |I j |

Q: What are we using?

A: Both. Cosine similarity, but limited to the cluster

Scoring

- Given a user, score each job and finally rank based on score:

    Score **= *w1*(Interaction score) + *w2*(no. of overlaps in user job-roles and item-titles) + *w3* (I[career level match]) + *w4* (I[discipline IDs match]) + *w5*(I[industry IDs match] + w6 (I[regions match])**

- Weights

Heuristic: w1 = 1, w2 = 100, w3 = 10, w4 = 12, w5 = 10, w6 = 5 & w7 = 2

- Linear Regression: Learn the weights

Evaluations

Approach	Score	Rank
Baseline	26,857.38	100
User User similarity	85,491.27	81
Interactions	180,112.15	72
Impressions	279,062.28	48
Heuristic Scoring	456,487.86	23
Linear Regression	468,767.08	20
Gradient Boosting	**473,758.10**	**20**

Conclusion